


Reproducing DeepVecFont-v2, and testing whether a sweep over it can resolve anything

Final project, part 2. Generative Models for Text and Images, Reichman University.

Paper: Wang, Wang, Yu, Zhu, Lian. *DeepVecFont-v2: Exploiting Transformers to Synthesize Vector Fonts with Higher Quality*. CVPR 2023. Course topic: autoregressive models. Code: fork of yizhiwang96/deepvecfont-v2. All our work is the main...repro diff.



The model reads panel 4 and writes panel 4. The reported metric compares panel 5. Everything this report measures happens in the gap between them.

The same letter B, from a font in the English test split. The model reads and writes the command sequence in panel 4. The score everyone reports compares panel 5. Most of what this report measures turns out to live in the distance between those two panels.

Abstract


We reproduce DeepVecFont-v2 on Chinese and English vector font synthesis, then run 21 single-factor changes to the architecture against it.

Our reconstruction scores 0.1621 on Chinese against the paper's reported 0.080. To find out why, we ran the authors' own released weights through our evaluation code. They score 0.1629, which is 0.0008 from our number. The reproduction is therefore faithful, and the distance to the published figure comes from the evaluation, not from our training.

Before changing anything we trained the baseline three times with different random seeds and measured how far apart the results landed. That spread is 0.0097. All 26 candidate configurations together span 0.0101. Changing the architecture moved the metric about as much as changing the seed. One change improves Chinese at every seed on both metrics, adding noise to the encoder during training. It does not carry over to English. One change reliably makes things worse. The rest are indistinguishable from noise.

1. Original architecture

DeepVecFont-v2 generates a full vector font from a few reference glyphs. It uses two modalities at once: a 64×64 raster image of each glyph and the glyph's outline as a sequence of drawing commands.



The vocabulary, on real data. A **glyph** is one character drawn in one style; a **font** is a whole alphabet drawn consistently. Its **outline** is built from segments, either straight or **cubic Bézier**, where each curve is fixed by two **on-curve points** it passes through and two **control points** it does not. To **render** is to turn that outline into pixels. The model never emits a raw coordinate: it picks one of 128 **bins** per axis, so it can only place a point on that grid. And it works **few-shot**, seeing four reference glyphs and having to draw the other forty-eight.

Encoders. A CNN encodes the reference images. A Transformer encodes the reference command sequences. The two are fused into one latent vector per font.

Decoders. Three heads read that latent. A CNN decodes a raster image. An autoregressive Transformer decoder emits the command sequence one token at a time. A second Transformer decoder then refines the whole sequence in parallel. **The refined sequence is what gets scored**, which matters later.

Representation. Each command is a type plus 8 coordinate arguments. Coordinates are quantized to a fixed grid of bins and predicted as a classification over bins, which is what makes the model autoregressive rather than a regressor. The paper adds a relaxation representation so gradients can flow through the discrete arguments, and a Bézier alignment loss that pulls sampled points on the predicted curve toward the target curve.

Loss. A weighted sum of an image L1 term, a command-type cross-entropy, a coordinate-bin cross-entropy, the Bézier alignment term, and a KL term on the latent.

Key hyperparameters, as shipped and as used by us:

	Value
Optimizer, learning rate	Adam, 2e-4
Batch size	32
Image size	64×64
Coordinate bins	128
Transformer width, encoder depth	512, 6
Reference glyphs	4 (English), 8 (Chinese)
Max sequence length	51 (English), 71 (Chinese)
Loss weights	image L1 10, Bézier 0.01, point 0.01
Parameters	136,055,207

1.1 Where the released code differs from the paper

We read the paper against the code. Four differences matter, and three of them became experiments in section 4.

Paper says	Code does
Coordinate arguments in 256 dimensions	quantizes to 128 bins
Self-refinement is a 2-layer decoder	builds 1 layer
Bézier loss weighted at 1.0	loss_w_aux = 0.01, 100× lower
The $N(0, I)$ perturbation is an inference-time device	applied in training, validation and test alike

Two more, found late, affect how our numbers compare to the published ones. The paper uses 10 samples per glyph on English and 50 on Chinese; we used 50 on both, which flatters our English figures. And the paper enlarges the Chinese training set ten times by affine augmentation, while our built dataset holds 1,272 entries over 212 base fonts, which is six times.

One more thing we found is not a paper discrepancy but is worth recording. The encoder builds cross-attention modules and never calls them. They amount to 23,204,402 parameters, 17% of the model, trained by the optimizer and used by nothing.

2. Paper results

The paper reports **reconstruction error**, written "Error" in its tables. The predicted outline is rendered to a 64×64 image, the ground truth glyph is rendered the same size, and the metric is the mean absolute difference between the two images. It is measured per glyph and averaged. Lower is better. The model draws several candidate outlines per glyph and the one with the highest overlap against the target is kept, so the reported number is a best-of-N score rather than a single draw.

Model	Error, English	Error, Chinese
DeepSVG	0.125	0.167
DeepVecFont	0.056	0.086
DeepVecFont-v2	0.052	0.080

Note the scale. The paper's whole improvement over its own predecessor on Chinese is 0.006. Every change we test in section 4 produces a difference around that size or smaller, and section 5 is mostly about whether such differences can be measured at all.

The paper also reports two ablations. Its English ablation spans 0.0069 in total across four rows. Its sampling-point study spans 0.0037 across six settings.


The metric from the course material. The brief asks for one metric studied in class alongside the paper's own. We use **SSIM**, which compares local structure rather than per-pixel difference and is standard for image quality. We implemented it in `eval_reconstruction_error.py` and checked it against skimage to 1e-16. We also report **s-IoU**, the intersection over union of the two rendered shapes, because the vector font literature uses it and because it turned out to disagree with L1 often enough to be worth carrying.

3. Reconstruction results

All numbers below are over the same 34 test fonts, best-of-50 decoding, with the ground truth taken from the dataset's pre-rendered images.

	Error (L1) ↓	s-IoU ↑	SSIM ↑
Paper, Chinese	0.080	not reported	not reported
Released weights, Chinese, epoch 600	0.1629	0.3225	0.4373
Our reconstruction, Chinese , 3-seed mean, 150 epochs	0.1621	0.2681	0.4425
Paper, English	0.052	not reported	not reported
Released weights, English, epoch 600	0.0658	0.7029	0.7181
Our reconstruction, English , 3-seed mean, 630 epochs	0.0597	0.7309	0.7374

Our reconstruction matches the released weights; neither reaches the published number




We did not reach the published numbers. We did reproduce the released model, and those are different claims.

On Chinese our reconstruction sits 0.0008 from the authors' own checkpoint, on a quarter of the training budget. On English ours scores better than all three released checkpoints. So the distance to 0.080 is not something our training did wrong. We ran the authors' weights through our code and got the same distance they would have.

Training longer does not close it. This was our leading explanation for three days. The released checkpoint has four times our budget and lands in the same place. We also trained our own Chinese runs to 600 epochs; their best checkpoints by validation loss come out at epochs 150 to 200, and they score 0.1583.

Part of the gap is the metric itself. The predicted outline is rendered with one rasterizer, and the ground truth image in the dataset was rendered with a different one. Two rasterizers, one comparison. If we render the ground truth outline through the same rasterizer as the prediction, the same fixed checkpoint moves from 0.1629 to 0.1174.



That single unstated choice is worth 0.0455 on Chinese, against a paper whose entire margin over its predecessor is 0.006. It is worth 0.0074 on English. It tracks the dataset rather than the model, so it is not a constant anyone can subtract once and forget.

A 0.16 Chinese Error is an average, and an average hides what the tail looks like. Section 5.4 shows both tails from the same six freshly-decoded test fonts, picked by L1 rather than by eye, alongside the qualitative baseline-versus-E9 comparison they belong next to.

What is still open. After accounting for the rasterizer, a Chinese-specific gap of about 0.037 remains, and we do not have an explanation we can prove. The best lead is the augmentation shortfall in section 1.1: our Chinese training set is six times augmented where the paper says ten. That would produce a Chinese-only gap while English, which needs no augmentation, reproduces cleanly. Testing it needs a retrain we did not have time for.

Every English figure here used 50 samples per glyph where the paper uses 10. Scoring one fixed checkpoint both ways puts that difference at 0.0027, so our English column is optimistic by roughly that much.

4. Improved architecture

4.1 What we did first, and why

The paper's own margin is 0.006. Before trying to beat it we checked whether we could measure anything that small. We trained the unmodified baseline three times, changing only the random seed, and looked at how far apart the three scores landed.

They span 0.0097 on Chinese and 0.0038 on English. That spread is the bar. Any change producing a smaller difference than this cannot be told apart from a lucky seed.

We then decomposed it. Re-decoding the same checkpoint accounts for 0.0011 of the Chinese figure and re-training accounts for the other 0.0082. Sampling more candidates per glyph therefore does not narrow the bar; only more seeds do.

This decision, made on day one, is the reason the rest of the report is worth reading. Two rules followed from it and were written down before any candidate ran:

- **A result must show the same sign at all three seeds.** One good run is not a result.
- **Every rule is fixed before the numbers arrive.** Where we later overruled one, we recorded the overrule and the reason rather than quietly dropping it.

4.2 The changes we tested

We ran 21 distinct single-factor changes. Each moves one thing and holds everything else at the baseline value. Every new option defaults to the released behaviour, and a 196-check

script asserts that, so a candidate’s default reproduces the baseline exactly.

Three candidates come from the paper-versus-code differences in section 1.1, and those are the ones we expected most from, because each asks whether the paper’s stated design or the shipped code is better.

Category from the brief	What we changed
Add normalization layers	terminal LayerNorm on the encoder (E1); batch norm in the image branch (E2)
Change the encoder or decoder	encoder width (E4); encoder depth 6→8 (E16); decoder feed-forward width (E17)
Change the latent dimension	bottleneck 512→256 and →1024 (E5)
Add residual or attention layers	refinement decoder depth 1→2 and →3 (E3); encoder depth (E16)
Modify the loss function	Bézier weight 0.01→0.1/0.3/1.0 (E7); ordinal label smoothing (E8); KL divergence (E9)
Change the noise schedule	encoder noise σ at train time (E9); warmup-cosine learning rate (E14)
Add regularization	dropout (E10); AdamW (E11); weight EMA (E15)

All seven categories in the brief are covered, and one representative of each was rerun at three seeds. That last part is the point: a coverage table built from single runs would show breadth and prove nothing, for the reason section 5.2 gives.

4.3 The main change: encoder noise

E9 sets the training-time encoder noise to $\sigma = 0.5$, half the released value.

The reason is in section 1.1. The code adds `x + torch.randn_like(x)` to the encoder output on every forward pass. The paper describes this perturbation as an inference-time device that simulates the uncertainty of human design. It is not described as a training regularizer, and it is applied as one.

We expected the released setting to be too strong. Full-strength noise on every training step should blur the latent more than the reconstruction task wants, while keeping it at test time still gives the sampling variety the paper is after. So we split the setting into two knobs, one for training and one for test, and swept the training one down.

The noise is still drawn even when scaled to zero, so the random number stream stays aligned across the sweep. Otherwise a change in σ would also be a change in initialization.

Two candidates did not survive contact with the code, and we report them rather than dropping them. Instance normalization cannot compute a per-instance variance at the model’s [32, 1024, 1, 1] bottleneck, so E2 ran with batch norm only. And we never ran an ablation of the 17% dead parameters, because removing them shifts the random stream for every module built afterwards, which would make it partly an initialization change.


5. Improved results

5.1 The comparison table

	Paper	Released weights	Our reconstruction	Our improved model (E9)
Chinese , Error (L1) ↓	0.080	0.1629	0.1621	0.1581
Chinese , s-IoU ↑	—	0.3225	0.2681	0.2952
Chinese , SSIM ↑	—	0.4373	0.4425	0.4479
English , Error (L1) ↓	0.052	0.0658	0.0597	0.0601
English , s-IoU ↑	—	0.7029	0.7309	0.7305
English , SSIM ↑	—	0.7181	0.7374	0.7358

Our columns are three-seed means. E9 improves all three Chinese metrics. It does not change English.

5.2 What the sweep returned



Twenty-six candidate configurations span 0.0101. Three runs of the unmodified model span 0.0097. Twenty-six draws from one distribution should cover roughly 2.3 times the range of three; here the ratio is 1.07.

This has a direct consequence we can measure on our own results. Compared against a single baseline run, **22 of 26 candidates beat the baseline**. Compared against the mean of three, **6 of 26** do. The only difference is that the single run happened to land 0.0048 above the mean, which is an ordinary draw. Twenty-two out of twenty-six reads as *almost anything helps*, and it would have been our headline if we had not measured the spread first.

We can also price the mistake. Ranking candidates against that single run and re-running the top three at two more seeds, **one of three survived**.

5.3 E9 in detail

Chinese, best-of-50, each candidate seed paired against the baseline seed of the same number.


Seed	E9 Error	Δ vs paired baseline	E9 s-IoU	Δ
1111	0.1588	-0.0074	0.2870	+0.0325
2222	0.1531	-0.0038	0.3170	+0.0454
3333	0.1623	-0.0009	0.2816	+0.0035
Mean	0.1581	-0.0040	0.2952	+0.0271

Neither mean clears the bar we set. The L1 gain of 0.0040 is inside the 0.0097 floor, and the s-IoU gain of 0.0271 is inside the 0.0315 floor. What E9 has is the other thing we required in advance: the same sign at every seed on both metrics, six out of six, on two metrics that barely correlate with each other ($r = -0.335$). If they were independent, six agreeing signs would happen by chance about 1 time in 64.

So E9 is a consistent direction, not a demonstrated gain of a stated size. We report it that way. A narrower floor is available from the higher sampling budget, 0.0236, and E9 would clear that one, but the bar was set at 0.0315 before anything ran and moving it afterwards would throw away the only thing that makes the sweep worth reading.

5.4 It does not transfer to English

The one improvement we found does not survive a change of script
shaded band is the seed-noise floor




We repeated E9 on English at three seeds, paired the same way. One seed favours it on both metrics and two go against it. Every difference is inside the English floor.


This was written down as a possible outcome before the runs started, so it is a result rather than a failed experiment: **the one improvement we found does not survive a change of script.**

What baseline and E9 actually draw, ground truth alongside, same font and same four characters on both languages:

Chinese · font 0000 · mean L1 over these 6 fonts -- baseline 0.133, +E9 0.131



English · font 0000 · mean L1 over these 6 fonts -- baseline 0.064, +E9 0.066




Same font, same four characters, both languages. Not a claim: section 5.3's reading is that E9's Chinese mean does not clear its own noise floor, and a picture cannot upgrade that -- this is what the numbers in section 5 look like as glyphs, nothing more.

English is close to solved at this scale; Chinese is not, and the two E9 columns are close enough to each other that a reader should not expect to see the improvement in section 5.3 by eye: 0.0040 L1 is about 16 of the image's 4096 pixels, a difference the metric can pick up averaged over hundreds of glyphs but not one a single rendered comparison reliably shows.

(The L1 figures printed above these panels use the ground-truth outline through our own rasterizer, section 3's "svg" convention, over only the six fonts decoded for this figure — not the raster-convention, 34-font numbers in sections 3 and 5.1, which they are not meant to match.)

The same six decoded Chinese fonts also bound the distribution behind that 0.163 mean. Section 3 said an average hides the tail; here is both ends of it, picked by L1 and not by eye, same models as above. First the low end:


Best Chinese baseline glyphs by L1, out of the 6 decoded test fonts (chosen by the number, not by eye)



The other tail of the same distribution as the failure strip below: these are what the low end of a 0.163 Chinese mean looks like, picked by the identical rule -- lowest L1 among the same six decoded fonts, not the most flattering glyphs found by scanning. Simple strokes with little overlap area reconstruct close to the ground truth; the failure strip shows what the model does with denser, more overlapping strokes instead.

Then the high end:

Worst Chinese baseline glyphs by L1, out of the 6 decoded test fonts (chosen by the number, not by eye)



The reported Chinese Error is 0.163 averaged over 34 fonts; these are what 0.16-ish actually looks like at the tail, not the average case. Section 3's point that the metric is floor-dominated on Chinese (0.142 of any score is cross-rasterizer disagreement, not model error) applies here too.

The pattern at the bad end is consistent: three of these four glyphs, all simple rectangular strokes in the ground truth, collapse into a solid filled block rather than an outline. E9 does not repeat that pattern on the same four glyphs — it fills one, leaves one nearly blank, and draws thin partial outlines on the other two — a different failure, not a smaller version of the same one. At the good end the pattern is close to the mirror image: what survives is glyphs with little stroke overlap to begin with, single bars and simple crosses, which is a property of the glyph rather than something either model did well. Neither tail is the rasterizer floor from section 3 (every comparison above renders the *same* outline on both sides; this is the model drawing the right or wrong thing), and the spread between them is the kind of thing a 64×64 pixel-disagreement count is well suited to catching.

5.5 The rest of the sweep

Everything else was null. Below are the one-per-category representatives, three seeds each on Chinese, at matched checkpoint epoch, against a floor of 0.0097 and 0.0315.

Change	Mean Δ Error	Mean Δ s-IoU	Reading
Batch norm in the image branch (E2)	−0.0066	+0.0221	same sign, under floor
Encoder width (E4)	+0.0059	+0.0173	mixed sign
Latent 256 (E5)	−0.0010	+0.0203	mixed sign
Bézier weight 0.1 (E7)	−0.0010	−0.0053	mixed sign
AdamW (E11)	−0.0011	+0.0017	mixed sign
Refinement depth 2 (E3)	−0.0004	+0.0007	mixed sign
Encoder depth 8 (E16)	+0.0025	−0.0045	mixed sign
Decoder width 2048 (E17)	+0.0005	+0.0051	mixed sign

E2 is the only one with a consistent sign on both metrics, and both of its means sit under the floor, so it does not pass either.

One change did clear the bar, and it made things worse. Adding a terminal LayerNorm to the encoder (E1) degrades Chinese s-IoU at all three seeds by a mean of 0.0376, which is 1.2 times the floor. This is the only effect anywhere in the project that exceeds its own bar in either direction.

That reading needed a correction, described in section 6.3. Its first version was -0.0760 , but the two worst runs had been scored at earlier checkpoints than the baseline they were compared against. Retraining and re-reading them at a matched epoch halves the effect. It survives.

Running the same changes on English disagreed with Chinese. Refinement depth 2 (E3) is mixed-sign and null on Chinese, and on English it clears both floors in the degrading direction at all three seeds. Latent 256 (E5) does the same. The same one-factor edit points in opposite directions on the two scripts the paper reports.

6. Discussion

6.1 What worked

The reproduction worked, and the way we checked it is the part worth keeping. Scoring the authors' released weights through our own code turned "we fell short of 0.080" into two separate and answerable questions: does our training match theirs, and does our evaluation match the paper's. The first is yes. The second is no, and we could measure part of why. Without a released checkpoint the honest conclusion available to us would have been the much weaker "probably undertrained", which earlier drafts of this report did in fact carry.

Measuring the seed spread before running any candidate also worked, and it changed every conclusion downstream. It cost three training runs on day one.

E9 worked in the limited sense described in section 5.3. Of our three paper-versus-code candidates, two produced our two clearest results, which suggests reading a paper against its implementation is a better source of hypotheses than guessing at hyperparameters.

6.2 What did not work

Eighteen of twenty-one changes did nothing measurable. We do not think this is because the changes were badly chosen. We think the benchmark cannot resolve changes of this size.

Our Chinese floor is 0.0097 and our English floor is 0.0038. Set those against the paper's own tables: its English ablation spans 0.0069 across four rows, with individual steps of 0.0031, 0.0028 and 0.0010. Its sampling-point study spans 0.0037 across six settings, with steps between 0.0002 and 0.0014. **Most rows in the paper's own ablations differ by less than our measured noise floor**, on the same metric and the same data.

This is not a claim that those results are wrong. It is a claim about what a single training run per configuration can support. Neither the paper nor the work it compares against reports a seed spread, so there is no way to tell from the published tables.

E9 not transferring to English is the other thing that did not work, and it is informative. So is the direction disagreement in section 5.5. Together they say that a one-language ablation licenses less than it looks like it does.

6.3 What we learned

A metric is a pipeline, not a formula. "Mean absolute error between two 64×64 images" is not enough information to reproduce a number. Two choices sit underneath it and neither is stated in the paper. Which rasterizer draws the ground truth is worth 0.0455 on Chinese. How many candidate outlines the best-of-N selection draws from is worth 0.0027 on English, and that quantity is pure sampling budget with no model in it. Both are larger than most of what we set out to measure.

A single-seed baseline reports the draw, not the change. Section 5.2 has the numbers: 22 of 26 versus 6 of 26, and a 1-in-3 survival rate for single-run leaders. This is the result we would most want someone else to check, and checking it is cheap. Take any published vector font ablation, retrain two rows at three seeds, and see whether the order holds.

Checkpoint selection was measuring the wrong thing. Checkpoints are chosen by a validation score, and no term in that score is the quantity being reported. Its largest term belongs to the image decoder, while what gets scored comes from the refinement decoder’s outline, rasterized. We measured the consequence four ways. Across candidates it ranks at Spearman $\rho = 0.125$. Within a run it puts epoch 100 above epoch 125 while every rendered metric says the opposite. It prefers the authors’ epoch-600 English checkpoint while every rendered metric prefers 500. And when we adopted it as the selection rule across five candidates, it never improved a rendered score and flipped two of five from null to degrading.

This cost us something real. E1’s original -0.0760 was measured on checkpoints the criterion had picked at epochs 100 and 125 against baselines at 150, so E1’s two worst runs were also its two earliest. At a matched epoch the effect halves to -0.0376 . It survived, but our one floor-clearing result was half an artifact of a rule nobody had audited.

Human judgement failed the same way. Two English runs were read from their training curves. One was stopped early on the shape of its curve and has no result. The other was read as going well and came back as our confirmed degrading result.

We did not switch the whole project to rendered-metric selection. Doing it properly means decoding and rasterizing the validation set at every checkpoint, and it would invalidate all 113 scored rows with four days left. We report it as measured and name the fix below.

6.4 Limitations

Every floor here rests on three seeds, so the bars themselves are uncertain; our Chinese s-IoU floor moved from 0.0401 to 0.0315 when we re-measured it. Checkpoint selection stayed on the criterion we just described as defective. English was scored on 34 fonts rather than the full 1,386, with the resulting bias measured at 0.0074 rather than assumed, and at 50 samples per glyph rather than the paper’s 10. Nothing was run at more than three seeds.

6.5 What we would do next

Rebuild the Chinese training set at the paper’s stated ten-times augmentation and retrain. It is cheap, about an hour per run, and it is the only open lead on the 0.037 in section 3.

Add a validation hook that decodes at one sample and rasterizes a small held-out set every few epochs. Checkpoint selection and scoring would then measure the same kind of thing, at a small fraction of the cost of a full evaluation.

Then reconsider the output head. The decoder is autoregressive over quantized coordinate bins. We measured what the quantization itself costs by encoding the ground truth outlines into the model’s own representation and decoding them back with no model involved: 0.1422 at unlimited precision and 0.1443 at the released 128-bin grid. The bins cost only 0.0021, so the grid is not the bottleneck, but the sequential factorization over it may be. A flow-matching head over continuous coordinates would remove the quantization and the sequential decode together. We designed and costed one and dropped it on schedule grounds rather than on merit; it is in `archive/FLOW_MATCHING_PLAN.md`.

7. References

1. Wang, Y., Wang, Y., Yu, L., Zhu, Y., Lian, Z. *DeepVecFont-v2: Exploiting Transformers to Synthesize Vector Fonts with Higher Quality*. CVPR 2023. arXiv:2303.14585.

2. Wang, Y., Lian, Z. *DeepVecFont: Synthesizing High-Quality Vector Fonts via Dual-Modality Learning*. ACM TOG 2021.
3. Carlier, A., Danelljan, M., Alahi, A., Timofte, R. *DeepSVG: A Hierarchical Generative Network for Vector Graphics Animation*. NeurIPS 2020.
4. Lopes, R. G., Ha, D., Eck, D., Shlens, J. *A Learned Representation for Scalable Vector Graphics*. ICCV 2019.
5. Liu, Y., Guo, F., Wang, Z., Zhang, D. *DualVector: Unsupervised Vector Font Synthesis with Dual-Part Representation*. CVPR 2023. Source of the SSIM / L1 / s-IoU reporting convention.
6. Wang, Z., Bovik, A. C., Sheikh, H. R., Simoncelli, E. P. *Image Quality Assessment: From Error Visibility to Structural Similarity*. IEEE TIP 2004. SSIM.
7. Xiong, R., et al. *On Layer Normalization in the Transformer Architecture*. ICML 2020. For E1.
8. Szegedy, C., Vanhoucke, V., Ioffe, S., Shlens, J., Wojna, Z. *Rethinking the Inception Architecture for Computer Vision*. CVPR 2016. Label smoothing, for E8.
9. Loshchilov, I., Hutter, F. *Decoupled Weight Decay Regularization*. ICLR 2019. AdamW, for E11.
10. Higgins, I., et al. *β -VAE: Learning Basic Visual Concepts with a Constrained Variational Framework*. ICLR 2017. For E12.
11. Jaegle, A., et al. *Perceiver: General Perception with Iterative Attention*. ICML 2021. The architecture the encoder’s configuration advertises and does not instantiate.
12. Wilcoxon, F. *Individual Comparisons by Ranking Methods*. Biometrics Bulletin 1945. The paired test used per seed.
13. Lipman, Y., et al. *Flow Matching Guide and Code*. arXiv:2412.06264. On the course reading list; basis for the head proposed in section 6.5.

Code and data. Upstream: github.com/yizhiwang96/deepvecfont-v2, mirrored on branch main; our work is the main...repro diff. Dataset: the authors’ released Chinese and English vector font sets, built by `data_utils/`. `RESULTS.csv` holds all 113 scored checkpoints, one row each. Every figure is regenerated from it by `report/make_figures.py` and every number in this report is checked against it by `report/verify_report.py`.